

A Mathematical Theory of Market World Models: Hierarchical Mean-Field Dynamics, Dual Stochastic Decomposition, and Financial Event Operator Algebras

HongJin HE

The Hong Kong University of Science and Technology Stanford University (IHP)

hhjkh@stanford.edu

July 2026 (*Preprint — submitted to arXiv:q-fin.MF*)

Abstract

We develop a rigorous mathematical theory for *market world models*— probabilistic dynamical systems that learn and simulate the full evolutionary law of financial markets rather than extracting statistical correlations through factor models. Our framework rests on four interlocking contributions.

(i) Dual Stochastic Decomposition. We prove that market noise admits a canonical Lévy–Itô decomposition into two orthogonal components: a *physical noise* driven by a Brownian motion (capturing fundamental information uncertainty) and a *behavioral noise* driven by a pure-jump Lévy process (capturing sentiment and herding effects). We derive a Cramér–Rao-type lower bound showing that prediction variance is bounded below by a sum of two irreducible terms arising from each component—a market-structure analogue of the Heisenberg uncertainty principle.

(ii) Financial Event Operator Algebra. Every corporate or macroeconomic event (IPO, merger, interest rate shock, etc.) is modeled as an affine operator on the market state space. We classify all such operators by their algebraic mode and prove that the full event algebra carries the structure of a *topological groupoid*—a generalisation of a semigroup required because mergers and spin-offs alter the dimension of the state space.

(iii) Hierarchical Mean-Field Theory. Real financial markets exhibit simultaneous mean-field game (MFG) and mean-field control (MFC) structures at different scales: nations interact as MFG agents while firms within each nation face a mixed MFG/MFC problem regulated by a central authority. We formulate a

two-level McKean–Vlasov system and prove existence and uniqueness of a *hierarchical Nash equilibrium* under Lipschitz, boundedness, and Lasry–Lions monotonicity conditions.

(iv) Stochastic Stability. Using the stochastic Lyapunov method for systems driven by the combined Brownian–Lévy noise, we prove that the market world model admits a unique invariant measure and establish sharp exponential convergence rates. Market bubbles and crises correspond precisely to local violations of the Lyapunov condition.

(v) Implementation Pathway. We bridge theory and practice by deriving a concrete computational realization: bipower-variation calibration of the dual noise parameters from high-frequency data (Stage 1); a Transformer VAE with predictive training objective for the Encoder E (Stage 2); a Deep Galerkin Method solver for the HJB equation coupled with neural fictitious play for the Fokker–Planck equation (Stage 3); and a gradient-based Controller that converts latent-space value-function gradients to risk-constrained portfolio weights (Stage 4). Convergence guarantees for each stage follow from the theoretical results in contributions (i)–(iv).

Taken together, these results constitute, to the best of our knowledge, the *first* systematic mathematical theory for market world models—one that is simultaneously axiomatically rigorous and computationally realizable: a self-contained, axiomatically grounded framework that connects the world-model paradigm of machine learning to the rigorous traditions of stochastic analysis and mean-field game theory in quantitative finance.

Keywords: world model, mean-field game, stochastic differential equation, Lévy process, groupoid, Lyapunov stability, market microstructure, prediction lower bound.

MSC 2020: 91G99, 60G51, 49N80, 22A22, 60H10.

Contents

1	Introduction	4
1.1	Motivation	4
1.2	Contributions	4
1.3	Related Work	5
2	Probabilistic Foundations	6
2.1	Probability Space and Filtration	6
2.2	Itô Stochastic Integral: The Correct Formulation	6
2.3	Lévy Processes and the Lévy–Khintchine Formula	7
2.4	Market State Space	7
3	The World Model as an Optimal Estimation Problem	8

4	Dual Stochastic Decomposition	8
4.1	Two Sources of Market Uncertainty	8
4.2	Canonical Decomposition Theorem	9
4.3	Prediction Lower Bound	10
5	Financial Event Operator Algebra	11
5.1	Affine Event Operators	11
5.2	Classification of Financial Events	11
5.3	Groupoid Structure of the Event Algebra	12
6	The E-Game-C Architecture	13
6.1	Limitations of the Original VMC World Model	13
6.2	Encoder Module E	13
6.3	Game Module: Mean-Field Equilibrium Operator	13
6.4	Controller Module C	14
7	Hierarchical Mean-Field Theory	15
7.1	Motivation: Multi-Scale Structure of Financial Markets	15
7.2	Two-Level McKean–Vlasov System	15
7.3	Hierarchical Nash Equilibrium	16
7.4	Assumptions	16
7.5	Main Theorem	17
8	Stochastic Lyapunov Stability	19
8.1	Lyapunov Stability under Dual Noise	19
8.2	Bubbles and Crises as Lyapunov Violations	20
9	Unified Evolution Equation	21
10	Implementation Pathway: From Theory to Market Prediction	21
10.1	Stage 1: Calibration of Dual Noise Parameters	22
10.2	Stage 2: The Encoder E — Variational Latent State Inference	22
10.3	Stage 3: Game Module — Approximate MFG Equilibrium in High Dimensions	23
10.4	Stage 4: Controller C — From Latent State to Trading Signals	24
10.5	End-to-End Training Pipeline	25
11	Discussion	26
11.1	Relation to Prior Work	26
11.2	Limitations and Future Work	26
12	Conclusion	27

1 Introduction

1.1 Motivation

The dominant paradigm in quantitative finance rests on two pillars: factor models, which represent asset returns as linear combinations of explanatory signals, and strategy models, which exploit persistent anomalies for profit (Fama, 1970; Ross, 1976). Despite their practical success, these approaches share a fundamental limitation: they learn *correlations* in historical data without learning the underlying *causal dynamics* of financial markets.

A market *world model* takes a different approach. Inspired by the world-model framework of Ha and Schmidhuber (2018) in reinforcement learning—where an agent maintains an internal model of its environment’s temporal dynamics—a market world model aims to learn the full evolutionary law governing market states, enabling counterfactual reasoning and scenario propagation unavailable to correlation-based methods.

Despite growing empirical interest (AWMLG, 2026), *no systematic mathematical theory* has been established for market world models. Existing work remains either purely algorithmic (lacking theoretical guarantees) or limited to single-scale models (ignoring the hierarchical multi-agent structure fundamental to financial markets). This paper provides the missing foundation.

1.2 Contributions

Our five main contributions, each grounded in formal theorems or algorithmic derivations, are:

- (i) **Dual Stochastic Decomposition** (Section 4). Theorem 4.2 establishes the canonical Lévy–Itô decomposition of market noise into physical and behavioral components. Theorem 4.4 provides a prediction lower bound—an information-theoretic impossibility result showing that no estimator can eliminate both components of forecast uncertainty.
- (ii) **Financial Event Operator Algebra** (Section 5). Definition 5.1 introduces affine event operators. Theorem 5.5 proves that the full algebra of financial events constitutes a topological groupoid.
- (iii) **Hierarchical Mean-Field Equilibrium** (Section 7). Theorem 7.4 is the main result of the paper: existence and uniqueness of a hierarchical Nash equilibrium for the two-level McKean–Vlasov system with a Lasry–Lions monotonicity structure.

- (iv) **Stochastic Lyapunov Stability** (Section 8). Theorem 8.2 establishes existence of a unique invariant measure and exponential convergence; Proposition 8.5 identifies the analytic signature of market bubbles and crashes.
- (v) **Implementation Pathway** (Section 10). We derive a concrete four-stage computational realization of the E-Game-C architecture: bipower-variation calibration of dual noise parameters; a Transformer VAE Encoder with predictive coupling; a Deep Galerkin Method (DGM) solver for the high-dimensional HJB equation with neural fictitious play for the Fokker–Planck equation; and a gradient-based Controller producing risk-constrained portfolio weights. Convergence of each stage follows from the theoretical results of contributions (i)–(iv).

1.3 Related Work

Mean-field games in finance. The MFG framework of Lasry and Lions (2007) and Huang et al. (2006) has been applied to price impact (Carmona and Lacker, 2015), systemic risk (Carmona and Mouzouni, 2023), and optimal liquidation (Carmona and Delarue, 2018). Multi-level MFG has been studied in the context of inter-bank lending (Carmona and Sun, 2018) but without the hierarchical coupling structure we introduce here.

Originality statement. To the best of our knowledge, this paper is the *first* to: (a) formally define a market world model as a probability-measure-valued map on an infinite-dimensional state space; (b) prove a Cramér–Rao-type lower bound that identifies two structurally distinct and irreducible components of market prediction uncertainty; (c) model the full algebra of financial events as a topological groupoid, capturing the dimension-changing nature of mergers and spin-offs; (d) establish existence and uniqueness of a hierarchical Nash equilibrium in a two-level McKean–Vlasov system that simultaneously exhibits MFG and MFC structures at different scales. Prior work on world models in finance is either empirical without theory (AWMLG 2026) or models only single-scale dynamics (Carmona and Delarue 2018).

Jump processes in finance. Lévy-driven asset prices are classical (Cont and Tankov, 2004; Kou, 2002). The decomposition of price innovations into a diffusion and a jump component has been studied for estimation purposes (Ait-Sahalia and Todorov, 2009), but not as a *structural* duality between physical and behavioral uncertainty.

World models and reinforcement learning. Since Ha and Schmidhuber (2018), world models have been central to model-based RL (Hafner et al., 2020, 2023). Their application to finance is nascent: AWMLG (2026) show that a generative world model

improves portfolio optimization; JPMorgan (2019) apply GANs to order-book simulation. No prior work provides a unified mathematical theory.

Groupoids in physics and mathematics. Groupoids generalize groups by allowing morphisms between different objects. Their use in noncommutative geometry (Connes, 1994) and geometric mechanics (Weinstein, 1996) is well-established; their appearance in financial modeling is novel.

2 Probabilistic Foundations

2.1 Probability Space and Filtration

Throughout the paper, we fix a complete probability space $(\Omega, \mathcal{F}, \mathbb{P})$ equipped with a right-continuous, complete filtration $\mathbb{F} = (\mathcal{F}_t)_{t \in [0, T]}$ for a fixed horizon $T \in (0, \infty)$.

We work with two independent sources of randomness:

- A standard $\mathbb{R}^m$ -valued Brownian motion $W = (W_t)_{t \geq 0}$, and
- A Poisson random measure N^η on $[0, T] \times \mathbb{R}^k$ with (non-zero) σ -finite intensity measure $\lambda \otimes \nu^\eta$, where λ denotes Lebesgue measure. The compensated measure is $\tilde{N}^\eta(dt, dz) := N^\eta(dt, dz) - \nu^\eta(dz) dt$.

Both W and N^η are adapted to $\mathbb{F}$ and mutually independent.

2.2 Itô Stochastic Integral: The Correct Formulation

A critical structural observation distinguishes our framework from ordinary differential equation (ODE) based approaches:

Proposition 2.1 (Non-differentiability of Brownian paths). *With probability one, $t \mapsto W_t$ is nowhere differentiable. Consequently, the expression $\sigma_t \cdot (dW_t/dt)$ is meaningless pointwise, and the shorthand notation $dS_t = \mu_t dt + \sigma_t dW_t$ represents an integral equation, not a differential equation.*

Proof. For any t , the increment $h^{-1}(W_{t+h} - W_t) \sim \mathcal{N}(0, h^{-1})$. As $h \rightarrow 0$, the variance $h^{-1} \rightarrow \infty$, so no L^2 limit exists. By the Paley–Wiener–Zygmund theorem, the set of ω for which $W(\omega)$ is differentiable at any single point has measure zero; a countable union argument extends this to all of $[0, T]$. $\square$

The precise meaning of the market evolution equation is therefore:

Definition 2.2 (Itô integral form). The market state process satisfies the *stochastic integral equation*:

$$\mathbf{S}_t = \mathbf{S}_0 + \int_0^t \mu(\mathbf{S}_u, u) du + \int_0^t \sigma_\tau(\mathbf{S}_u) dW_u + \int_0^t \int_{\mathbb{R}^k} z \tilde{N}^\eta(du, dz) + \sum_{\tau_w \leq t} (T_w - I) \mathbf{S}_{\tau_w^-}, \quad (1)$$

where the second and third Itô integrals are defined as $L^2(\Omega)$ limits of simple predictable approximants, and the last sum runs over all financial event times $\tau_w \in [0, t]$.

The following isometry is fundamental:

Proposition 2.3 (Itô isometry). *For any $\mathbb{F}$ -predictable process σ with $\mathbb{E}[\int_0^T \|\sigma_t\|^2 dt] < \infty$,*

$$\mathbb{E} \left[\left\| \int_0^t \sigma_u dW_u \right\|^2 \right] = \mathbb{E} \left[\int_0^t \|\sigma_u\|^2 du \right].$$

The adaptedness requirement $\sigma_u \in \mathcal{F}_u$ encodes the economic constraint of no look-ahead: trading signals may not depend on future information.

2.3 Lévy Processes and the Lévy–Khintchine Formula

Definition 2.4 (Lévy process). A process $L = (L_t)_{t \geq 0}$ is a *Lévy process* if (i) $L_0 = 0$ a.s., (ii) L has independent, stationary increments, and (iii) L is continuous in probability.

Theorem 2.5 (Lévy–Khintchine, Sato 1999). *The characteristic function of L_t has the form $\mathbb{E}[e^{i\xi^\top L_t}] = e^{-t\Psi(\xi)}$, where the Lévy exponent is:*

$$\Psi(\xi) = -ib^\top \xi + \frac{1}{2} \xi^\top C \xi + \int_{\mathbb{R}^k} \left(1 - e^{i\xi^\top z} + i\xi^\top z \mathbb{1}_{|z| \leq 1} \right) \nu(dz),$$

with drift $b \in \mathbb{R}^k$, diffusion matrix $C \in \mathbb{R}^{k \times k}$ positive semi-definite, and Lévy measure ν satisfying $\int_{\mathbb{R}^k} (1 \wedge |z|^2) \nu(dz) < \infty$. The triple (b, C, ν) is the Lévy characteristic and determines L in law.

2.4 Market State Space

Definition 2.6 (Single-asset state vector). The state of a single asset at time t is the vector $s_t = (p_t, v_t, \ell_t, \kappa_t, \iota_t)^\top \in \mathbb{R}^5$, where:

Coord.	Symbol	Economic meaning	Domain
1	p_t	Log-price	$\mathbb{R}$
2	v_t	Log-trading volume	$\mathbb{R}_+$
3	ℓ_t	Leverage ratio	$[0, \infty)$
4	κ_t	Outstanding shares (log)	$\mathbb{R}_+$
5	ι_t	Public information disclosure level	$[0, 1]$

For n assets, the joint state is $\mathbf{S}_t \in \mathbb{R}^{5n}$. In the continuum limit ($n \rightarrow \infty$) relevant for theoretical analysis, we work in the separable Hilbert space $\mathcal{H} = L^2([0, 1] \times [0, T], \mathbb{R}^5)$ with inner product $\langle f, g \rangle = \int_0^1 \int_0^T f(u, t) \cdot g(u, t) dt du$.

3 The World Model as an Optimal Estimation Problem

Definition 3.1 (Market world model). A *market world model* with parameters θ is a measurable map

$$\mathcal{M}_\theta : \mathcal{F}_t \times (0, \infty) \rightarrow \mathcal{P}(\mathcal{H})$$

such that $\mathcal{M}_\theta(\mathcal{F}_t, h) \approx \mathbb{P}(\mathbf{S}_{t+h} \in \cdot \mid \mathcal{F}_t)$.

Theorem 3.2 (Prediction error decomposition). For any $\mathcal{F}_t$ -measurable estimator $\hat{\mathbf{S}}_{t+h}$ of $\mathbf{S}_{t+h}$,

$$\mathbb{E} \left[\left\| \hat{\mathbf{S}}_{t+h} - \mathbf{S}_{t+h} \right\|_{\mathcal{H}}^2 \right] = \|\text{Bias}_\theta\|^2 + \underbrace{\text{Var}_\tau(h)}_{\text{physical}} + \underbrace{\text{Var}_\eta(h)}_{\text{behavioral}}, \quad (2)$$

where the last two terms are strictly positive and model-independent (see Theorem 4.4).

Proof. The bias–variance decomposition is standard. For the physical variance, write $\mathbf{S}_{t+h} = \hat{\mathbf{S}}_{t+h} + (\mathbf{S}_{t+h} - \mathbb{E}[\mathbf{S}_{t+h} \mid \mathcal{F}_t]) + (\mathbb{E}[\mathbf{S}_{t+h} \mid \mathcal{F}_t] - \hat{\mathbf{S}}_{t+h})$ and expand using $\mathbb{E}[M_t(A_t - B)] = 0$ for a martingale M and $\mathcal{F}_t$ -measurable B . The decomposition into physical and behavioral components follows from the independence established in Theorem 4.2. $\square$

4 Dual Stochastic Decomposition

4.1 Two Sources of Market Uncertainty

Standard models (e.g., Black–Scholes) conflate two qualitatively different sources of market randomness into a single Brownian motion. We argue this conflation is structurally incorrect:

- **Physical noise** τ_t : Arises from fundamental information asymmetry, parameter estimation error, and market microstructure frictions. It accumulates *continuously* along trading time—analogueous to thermal noise in physical systems.
- **Behavioral noise** η_t : Arises from investor sentiment, herding, panic selling, and euphoric buying. It manifests as *discontinuous jumps*—sudden regime shifts that cannot be generated by any continuous Brownian path.

Assumption 4.1 (Dual noise structure). The noise in equation (1) decomposes as:

$$\text{Noise}_t = \underbrace{\sigma_\tau(\mathbf{S}_t) dW_t^{(\tau)}}_{\text{physical}} + \underbrace{dJ_t^{(\eta)}}_{\text{behavioral}}, \quad (3)$$

where $W^{(\tau)}$ is a standard Brownian motion and $J_t^{(\eta)} = \int_0^t \int_{\mathbb{R}^k} z \tilde{N}^\eta(ds, dz)$ is an independent pure-jump Lévy process with Lévy measure ν^η satisfying $\int_{|z|>1} |z|^2 \nu^\eta(dz) < \infty$ and $[W^{(\tau)}, J^{(\eta)}]_t = 0$ a.s.

4.2 Canonical Decomposition Theorem

Theorem 4.2 (Dual Lévy–Itô decomposition). *Under Assumption 4.1, the total noise process $L_t := \sigma_\tau W_t^{(\tau)} + J_t^{(\eta)}$ has a unique Lévy characteristic (b, C, ν) that decomposes as:*

$$b = b_\tau + b_\eta, \quad C = \Sigma_\tau := \sigma_\tau \sigma_\tau^\top, \quad \nu = \nu^\eta, \quad (4)$$

with Σ_τ (the diffusion matrix) and ν^η (the jump intensity measure) uniquely identified by the characteristic function:

$$\mathbb{E}\left[e^{i\xi^\top L_t}\right] = \exp\left(-t\left[\frac{1}{2}\xi^\top \Sigma_\tau \xi + \int_{\mathbb{R}^k} \left(1 - e^{i\xi^\top z} + i\xi^\top z \mathbb{1}_{|z|\leq 1}\right) \nu^\eta(dz)\right]\right). \quad (5)$$

In particular, $\Sigma_\tau \neq 0$ and $\nu^\eta \neq 0$ are orthogonal in the sense that no jump measure can generate a continuous Gaussian component.

Proof. By independence of $W^{(\tau)}$ and $J^{(\eta)}$, the characteristic function of L_t factors as

$$\mathbb{E}[e^{i\xi^\top L_t}] = \mathbb{E}[e^{i\xi^\top \sigma_\tau W_t}] \cdot \mathbb{E}[e^{i\xi^\top J_t^\eta}] = e^{-t\Psi_\tau(\xi)} \cdot e^{-t\Psi_\eta(\xi)},$$

where $\Psi_\tau(\xi) = \frac{1}{2}\xi^\top \Sigma_\tau \xi$ (Gaussian) and $\Psi_\eta(\xi) = \int (1 - e^{i\xi^\top z} + i\xi^\top z \mathbb{1}_{|z|\leq 1}) \nu^\eta(dz)$ (pure-jump Lévy). Adding the two Lévy exponents gives (5). Uniqueness of the Lévy characteristic (Theorem 2.5) and the canonical decomposition of Lévy processes (Sato, 1999) establish that Σ_τ and ν^η are uniquely determined. The orthogonality statement reflects the Lévy–Khintchine classification: the Gaussian part is determined by $C = \Sigma_\tau$ alone, the jump part by ν^η alone; no assignment of mass to ν^η at any point in $\mathbb{R}^k$ can contribute to C . $\square$

Remark 4.3 (Heavy tails of behavioral noise). Empirical evidence (Cont and Bouchaud, 2000; Lux and Marchesi, 1999) supports $\nu^\eta(|z| > x) \sim x^{-\alpha}$ for $\alpha \in (1, 2)$ (heavy-tailed but with finite variance), consistent with the power-law tails observed in financial return distributions and inconsistent with any Gaussian model.

4.3 Prediction Lower Bound

Theorem 4.4 (Cramér–Rao-type prediction lower bound). *Let $\mathbf{S}$ satisfy (1) under Assumption 4.1. For any $\mathcal{F}_t$ -measurable unbiased estimator $\hat{\mathbf{S}}_{t+h}$ of $\mathbf{S}_{t+h}$ and any horizon $h > 0$:*

$$\text{Var}\left(\hat{\mathbf{S}}_{t+h} - \mathbf{S}_{t+h}\right) \geq \underbrace{\sigma_\tau^2 h}_{\text{physical lower bound}} + \underbrace{\lambda_\eta m_2^\eta h}_{\text{behavioral lower bound}} =: \text{CR}(h), \quad (6)$$

where $\sigma_\tau^2 = \text{tr}(\Sigma_\tau)$, $\lambda_\eta > 0$ is the jump intensity, and $m_2^\eta = \int_{\mathbb{R}^k} |z|^2 \nu^\eta(dz)$.

Proof. Step 1 (Physical component). Conditional on $\mathcal{F}_t$, the increment $\mathbf{S}_{t+h}^{(\tau)} - \mathbf{S}_t^{(\tau)} := \int_t^{t+h} \sigma_\tau dW_u$ is Gaussian with mean zero and covariance $\Sigma_\tau h$. For any unbiased estimator $\hat{A}$ of a Gaussian parameter with Fisher information I , the standard Cramér–Rao inequality gives $\text{Var}(\hat{A}) \geq I^{-1}$. Here $I^{-1} = \Sigma_\tau h$, so $\text{Var}(\hat{\mathbf{S}}_{t+h}^{(\tau)}) \geq \text{tr}(\Sigma_\tau)h = \sigma_\tau^2 h$.

Step 2 (Behavioral component). Let $K_{t,t+h} = N^\eta([t, t+h] \times \mathbb{R}^k)$ be the number of jumps in $[t, t+h]$. Given $K_{t,t+h} = n$, the jump sizes are i.i.d. with common distribution $\nu^\eta(dz)/\lambda_\eta$, so the total jump contribution has variance $n \cdot m_2^\eta / \lambda_\eta$. Integrating over $K_{t,t+h} \sim \text{Poisson}(\lambda_\eta h)$:

$$\text{Var}\left(\int_t^{t+h} \int z \tilde{N}^\eta(du, dz)\right) = \lambda_\eta h \cdot m_2^\eta.$$

Since $\hat{\mathbf{S}}_{t+h}$ cannot predict the future jump realisation without observing N^η beyond time t , this variance is irreducible.

Step 3 (Independence). By Assumption 4.1, physical and behavioral increments are independent, so their variances add: $\text{Var}(\mathbf{S}_{t+h}^{(\tau)} + J_{t+h}^{(\eta)} - J_t^{(\eta)}) = \sigma_\tau^2 h + \lambda_\eta m_2^\eta h = \text{CR}(h)$. $\square$

Remark 4.5 (Analogy with quantum mechanics). Theorem 4.4 is the financial analogue of the Heisenberg uncertainty principle: even with perfect knowledge of $\mathcal{F}_t$ and an optimal model, the sum of prediction variances across physical and behavioral channels cannot fall below $\text{CR}(h)$. This bound is a structural feature of the market, not a limitation of methodology.

Corollary 4.6 (Time-varying temperature parameter). *The scalar temperature parameter τ_t of the world model (Ha and Schmidhuber, 2018) has the closed-form expression:*

$$\tau_t = \sqrt{\sigma_\tau^2(\mathbf{S}_t) + \lambda_\eta(t) m_2^\eta(t)},$$

and is a $\mathbb{F}$ -adapted, time-varying quantity determined by local noise intensities. Training the world model to match the observed local variance provides a principled, data-driven estimate of τ_t .

5 Financial Event Operator Algebra

5.1 Affine Event Operators

Definition 5.1 (Affine event operator). For each financial event type $w \in \mathcal{W}$, the associated *event operator* is the affine map:

$$T_w(s) = A_w s + b_w + \Sigma_w \varepsilon_w, \quad \varepsilon_w \sim \mathcal{N}(0, I), \quad (7)$$

where $A_w \in \mathbb{R}^{d' \times d}$ (with $d' = d$ or $d' \neq d$ depending on the event type), $b_w \in \mathbb{R}^{d'}$, and $\Sigma_w \in \mathbb{R}^{d' \times d'}$ encodes event-specific announcement uncertainty.

5.2 Classification of Financial Events

Financial events fall into three algebraic modes based on how they act on the state space:

Definition 5.2 (Three modes of action). (I) **Local endomorphism** ($d' = d$, acts on one asset): stock split, share repurchase, earnings announcement, dividend payment. $A_w \in \mathbb{R}^{d \times d}$; the state space dimension is preserved.

(II) **Global tensor action** (acts on all n assets via Kronecker product): central bank rate change, systemic shock, sector-wide regulation. $T_w^{\text{global}} = \Lambda_w \otimes I_d$, with $\Lambda_w \in \mathbb{R}^{n \times n}$ encoding heterogeneous sensitivities.

(III) **Pairwise morphism** ($d' \neq d$, changes entity count): merger & acquisition ($n \rightarrow n - 1$), spin-off / IPO ($n \rightarrow n + 1$). A_w is *non-square*; the state space dimension changes.

Proposition 5.3 (Information irreversibility). *For any Type-I event w with natural inverse w^{-1} (e.g., stock split followed by reverse split),*

$$T_{w^{-1}} \circ T_w = I + \mathcal{E}_w^{\text{info}},$$

where $\mathcal{E}_w^{\text{info}} \neq 0$ is an operator acting non-trivially on the information subspace $\{\iota\text{-coordinate}\}$ of the state vector. Even when A_w is invertible, the event releases information that cannot be unannounced: matrix invertibility does not imply physical reversibility.

Proof. Let $s = (p, v, \ell, \kappa, \iota)^\top$. For a stock split with ratio $r > 1$: $A_w = \text{diag}(1, 1, 1, r, 1)$, $b_w = 0$. The split also increases the disclosure level: $\iota \mapsto \min(\iota + \delta_w, 1)$ for some $\delta_w > 0$. Hence b_w has a positive entry in the ι -coordinate. The inverse operation $A_{w^{-1}} = \text{diag}(1, 1, 1, r^{-1}, 1)$ restores κ but leaves ι elevated, so $(T_{w^{-1}} \circ T_w)(s)_\iota = \iota + \delta_w \neq \iota$. $\square$

5.3 Groupoid Structure of the Event Algebra

Type-III events (mergers and spin-offs) change the state space dimension n , making the operator A_w non-square. This prevents the collection of all event operators from forming a semigroup (which requires a single carrier set). The correct algebraic structure is a *groupoid*.

Definition 5.4 (Financial event groupoid). Define the *financial event groupoid* $\mathcal{G}_{\text{fin}} = (\mathcal{W}, \circ, \text{src}, \text{tgt}, \text{id})$ by:

- **Objects:** $\text{Ob}(\mathcal{G}) = \{\mathcal{H}_S : S \text{ is a valid } n\text{-asset market configuration, } n \in \mathbb{N}\}.$
- **Morphisms:** For $T_w \in \text{Hom}(\mathcal{H}_S, \mathcal{H}_{S'})$, $\text{src}(T_w) = \mathcal{H}_S$ and $\text{tgt}(T_w) = \mathcal{H}_{S'}.$
- **Composition:** $T_{w_2} \circ T_{w_1}$ is defined if and only if $\text{src}(T_{w_2}) = \text{tgt}(T_{w_1}).$
- **Identities:** $\text{id}_{\mathcal{H}_S} = I_{d \cdot |S|}.$
- **Topology:** Endow each $\text{Hom}(\mathcal{H}_S, \mathcal{H}_{S'})$ with the operator norm topology; the whole space with the disjoint-union topology.

Theorem 5.5 (Groupoid structure). $\mathcal{G}_{\text{fin}}$ is a topological groupoid.

Proof. We verify the four groupoid axioms:

(i) *Well-defined composition.* Given $T_{w_1} : \mathcal{H}_S \rightarrow \mathcal{H}_{S'}$ and $T_{w_2} : \mathcal{H}_{S'} \rightarrow \mathcal{H}_{S''}$, their composition $T_{w_2} \circ T_{w_1} : \mathcal{H}_S \rightarrow \mathcal{H}_{S''}$ is the standard operator composition, which is well-defined in Banach spaces whenever the intermediate spaces match. The dimensionality compatibility condition $\text{src}(T_{w_2}) = \text{tgt}(T_{w_1})$ is exactly this matching requirement.

(ii) *Associativity.* Operator composition in Banach spaces is associative: $(T_{w_3} \circ T_{w_2}) \circ T_{w_1} = T_{w_3} \circ (T_{w_2} \circ T_{w_1})$ whenever both sides are defined.

(iii) *Identity morphisms.* For each $\mathcal{H}_S$, the identity operator $\text{id}_{\mathcal{H}_S} = I_{d|S|}$ satisfies $T_w \circ \text{id}_{\mathcal{H}_S} = T_w$ and $\text{id}_{\mathcal{H}_{S'}} \circ T_w = T_w$ for any $T_w \in \text{Hom}(\mathcal{H}_S, \mathcal{H}_{S'}).$

(iv) *Inverses (partial).* For Type-I and Type-II events with invertible A_w (i.e., $\det(A_w) \neq 0$), the inverse T_w^{-1} exists as an operator and $T_w \circ T_w^{-1} = \text{id}_{\mathcal{H}_S}$ (up to the information residual in Proposition 5.3). For Type-III events (non-square A_w), the Moore–Penrose pseudoinverse T_w^+ satisfies $T_w T_w^+ T_w = T_w$ but is not a full inverse—this is the natural partial-inverse structure of a groupoid.

(v) *Continuity.* Under the operator norm topology, $(T_{w_1}, T_{w_2}) \mapsto T_{w_2} \circ T_{w_1}$ is continuous (composition of bounded operators is jointly continuous in the norm topology). Similarly, $T_w \mapsto T_w^+$ is continuous on the set of operators with fixed rank. Hence $\mathcal{G}_{\text{fin}}$ is a topological groupoid. $\square$

Corollary 5.6 (Endomorphism monoid as a sub-structure). *The collection of Type-I events $\mathcal{W}_I \subset \mathcal{W}$ acting on a fixed n -asset space forms an endomorphism monoid with identity I_{5n} . For Type-I events with $\det(A_w) \neq 0$ this monoid is a group.*

6 The E-Game-C Architecture

6.1 Limitations of the Original VMC World Model

The V–M–C architecture of Ha and Schmidhuber (2018) consists of: a Variational AutoEncoder (**V**) that compresses high-dimensional observations to a latent code z_t ; a recurrent world model (**M**) that predicts z_{t+1} from (z_t, a_t) ; and a linear controller (**C**) trained by evolution strategies.

Direct application to finance faces three structural obstacles: (a) the input is heterogeneous (prices, news, event sequences), not pixels; (b) the “environment” is not fixed—it is the aggregate outcome of all agents’ decisions, creating a reflexive feedback loop; (c) the world model module M implicitly assumes the environment is a fixed dynamical system, ignoring the multi-agent strategic structure that is fundamental to financial markets.

6.2 Encoder Module E

Definition 6.1 (Information encoder). The encoder $\mathcal{E}_\phi : \mathcal{F}_t \rightarrow \mathcal{Z}$ maps the information set at time t to a latent code $z_t \in \mathcal{Z} \subset \mathbb{R}^{d_z}$ via a variational autoencoder (Kingma and Welling, 2014):

$$z_t = \mathcal{E}_\phi(\mathbf{x}_t), \quad \mathbf{x}_t = (p_t, v_t, \text{news}_t, \text{events}_t, \dots) \in \mathcal{F}_t,$$

trained by maximising the ELBO:

$$\mathcal{L}_{\text{ELBO}}(\phi, \psi) = \mathbb{E}_{z \sim q_\phi} [\log p_\psi(\mathbf{x}|z)] - D_{\text{KL}}(q_\phi(z|\mathbf{x}) \| p(z)).$$

6.3 Game Module: Mean-Field Equilibrium Operator

The core innovation of our architecture is replacing the recurrent world model M with a mean-field equilibrium operator.

Definition 6.2 (Market mean-field game). The market consists of a continuum of agents $u \in [0, 1]$. Agent u ’s latent state evolves as:

$$dz_t^u = f(z_t^u, a_t^u, \mu_t, u) dt + \sigma dW_t^u, \quad z_0^u \sim \rho_0, \quad (8)$$

where $a_t^u \in \mathcal{A}$ is the action, $\mu_t = \text{Law}(z_t^u)$ is the *mean field* (distribution of all agents), and $(W^u)_{u \in [0, 1]}$ are independent Brownian motions. Agent u minimises the cost:

$$J^u(a) = \mathbb{E} \left[\int_0^T \ell(z_t^u, a_t^u, \mu_t, u) dt + g(z_T^u, \mu_T, u) \right]. \quad (9)$$

Theorem 6.3 (MFG equilibrium existence). *Under the following conditions:*

(A1) f and ℓ are jointly continuous in all arguments and Lipschitz in (z, a, μ) with constant L_f ;

(A2) The action space $\mathcal{A}$ is compact and convex;

(A3) There exists a unique minimiser $\hat{a}(z, \mu, u) = \arg \min_{a \in \mathcal{A}} \mathcal{H}(z, a, \mu, u, p)$ for the Hamiltonian $\mathcal{H}$ with adjoint p ;

(A4) (Lasry–Lions monotonicity) $\int [\ell(z, \mu, a) - \ell(z, \mu', a)] d(\mu - \mu')(z) \geq 0$ for all μ, μ', a ,

there exists a unique mean-field game equilibrium $(\hat{a}^*, \hat{\mu})$ satisfying: (i) $\hat{a}_t^* = \arg \min_a J^u(a | \hat{\mu})$ for μ -a.e. u ; (ii) $\hat{\mu}_t = \text{Law}(z_t^u | \hat{a}^*)$ for all t .

Proof. Define the fixed-point map $\Phi : C([0, T]; \mathcal{P}_2(\mathcal{Z})) \rightarrow C([0, T]; \mathcal{P}_2(\mathcal{Z}))$ by: given a flow $\mathbf{m} = (\mu_t)$, solve the forward SDE (8) with optimal feedback $\hat{a}(z, \mu_t, u)$ and output $\Phi(\mathbf{m})_t = \text{Law}(z_t^u)$.

Self-mapping. By (A1)–(A2) and Gronwall, solutions of (8) satisfy a uniform L^2 bound, so $\Phi(\mathbf{m}) \in C([0, T]; \mathcal{P}_2(\mathcal{Z}))$.

Compactness. Tightness of $\{\Phi(\mathbf{m}_n)\}$ for bounded sequences $\{\mathbf{m}_n\}$ follows from Arzela–Ascoli applied to the SDE solutions.

Continuity. If $\mathbf{m}_n \rightarrow \mathbf{m}$ in $C([0, T]; \mathcal{P}_2(\mathcal{Z}))$ (Wasserstein-2 topology), then by (A1) and stability of SDEs, $\Phi(\mathbf{m}_n) \rightarrow \Phi(\mathbf{m})$.

Schauder’s theorem. Φ maps the closed convex set $\mathcal{K} = \{\mathbf{m} : \sup_t W_2(\mu_t, \rho_0)^2 \leq R\}$ (for R large enough) into itself, continuously, and $\mathcal{K}$ is compact. Schauder’s fixed-point theorem (Evans, 2010) yields a fixed point $\hat{\mathbf{m}}$ with $\Phi(\hat{\mathbf{m}}) = \hat{\mathbf{m}}$, i.e., a MFG equilibrium.

Uniqueness. The Lasry–Lions condition (A4) ensures that any two equilibria $(\hat{a}^*, \hat{\mu})$ and $(\hat{a}^{**}, \hat{\mu}^{**})$ satisfy $W_2(\hat{\mu}_t, \hat{\mu}_t^{**}) = 0$ for all t (by a standard energy estimate using the monotonicity of ℓ), hence uniqueness. $\square$

6.4 Controller Module C

Definition 6.4 (Optimal controller). The controller $\pi^* : \mathcal{Z} \times \mathcal{H}_{\text{RNN}} \rightarrow \mathcal{A}$ maps the current latent state z_t and world-model hidden state h_t to an optimal action:

$$\pi^*(z_t, h_t) = \arg \max_{a \in \mathcal{A}} Q^*(z_t, h_t, a),$$

where Q^* is the optimal action-value function satisfying the HJB equation:

$$0 = \partial_t Q + \sup_{a \in \mathcal{A}} [f(z, a, \hat{\mu}_t)^\top \nabla_z Q + \ell(z, a, \hat{\mu}_t)] + \frac{1}{2} \sigma^2 \Delta_z Q.$$

7 Hierarchical Mean-Field Theory

7.1 Motivation: Multi-Scale Structure of Financial Markets

A key structural observation motivates the hierarchical extension:

1. At the **national/regional scale**, sovereign entities (countries, central banks, regulatory bodies) interact without a supranational central planner—a MFG structure.
2. Within each nation, individual firms are subject to regulation and monetary policy set by the central bank—a partial MFC structure.
3. At the **firm level**, individual investors within a firm’s investor base again interact without coordination—a MFG structure.

This hierarchical interplay—MFG at the macro level, mixed MFG/MFC at the meso level—cannot be captured by a single-level mean-field model.

7.2 Two-Level McKean–Vlasov System

Let $N \in \mathbb{N}$ denote the number of groups (nations) and M the size of the continuum of agents within each group.

Level 1 (macro-group dynamics). Each group $j \in \{1, \dots, N\}$ has a representative state $\xi_t^j \in \mathbb{R}^{d_1}$ evolving as:

$$d\xi_t^j = b_1(\xi_t^j, \nu_t^{(1)}, \alpha_t^j) dt + \sigma_1 dB_t^j, \quad \xi_0^j = \xi_0^j \in \mathbb{R}^{d_1}, \quad (10)$$

where $\nu_t^{(1)} = \frac{1}{N} \sum_{k=1}^N \delta_{\xi_t^k}$ is the empirical measure of macro-group states, α_t^j is the policy (control) of group j , and $(B^j)_{j=1}^N$ are independent Brownian motions. Group j minimises the macro-cost:

$$J_1^j(\alpha^j) = \mathbb{E} \left[\int_0^T f_1(\xi_t^j, \nu_t^{(1)}, \alpha_t^j) dt + g_1(\xi_T^j, \nu_T^{(1)}) \right]. \quad (11)$$

Level 2 (micro-agent dynamics). Within group j , a continuum of agents $u \in [0, 1]$ has states $x_t^{j,u} \in \mathbb{R}^{d_2}$ evolving as:

$$dx_t^{j,u} = b_2(x_t^{j,u}, \mu_{j,t}^{(2)}, \xi_t^j, a_t^{j,u}) dt + \sigma_2 dW_t^{j,u}, \quad (12)$$

where $\mu_{j,t}^{(2)} = \int_0^1 \delta_{x_t^{j,u}} du$ is the intra-group mean field, $a_t^{j,u}$ is agent u ’s control, and $(W^{j,u})$ are i.i.d. Each agent minimises:

$$J_2^{j,u}(a) = \mathbb{E} \left[\int_0^T f_2(x_t^{j,u}, \mu_{j,t}^{(2)}, \xi_t^j, a_t^{j,u}) dt + g_2(x_T^{j,u}, \mu_{j,T}^{(2)}, \xi_T^j) \right]. \quad (13)$$

Coupling condition. The Level-1 and Level-2 dynamics are coupled through the *aggregation functional*

$$\Psi_j(\mu_{j,\cdot}^{(2)}) := \int_{\mathbb{R}^{d_2}} \phi(x) \mu_{j,t}^{(2)}(dx), \quad t \in [0, T], \quad (14)$$

for a bounded Lipschitz test function ϕ ; the macro-group cost J_1^j depends on $\Psi_j(\mu_{j,\cdot}^{(2)})$ through the drift b_1 .

7.3 Hierarchical Nash Equilibrium

Definition 7.1 (Hierarchical Nash equilibrium). A pair $((\alpha^{j,*})_{j=1}^N, (a^{j,u,*})_{j,u})$ is a *hierarchical Nash equilibrium* if:

- (i) For each group j and the resulting Level-2 equilibrium $\hat{\mu}_j^{(2)}[\xi^j]$ (defined in Definition 7.5 below), the macro policy satisfies:

$$J_1^j(\alpha^{j,*}) \leq J_1^j(\alpha^j) \quad \forall \alpha^j, \text{ given } (\alpha^{k,*})_{k \neq j};$$

- (ii) For each group j and each agent $u \in [0, 1]$, given ξ^j and $\mu_{j,t}^{(2)} = \hat{\mu}_{j,t}^{(2)}[\xi^j]$:

$$J_2^{j,u}(a^{j,u,*}) \leq J_2^{j,u}(a^{j,u}) \quad \forall a^{j,u}.$$

7.4 Assumptions

Assumption 7.2 (Regularity at Level 2). (L1) b_2, f_2, g_2 are jointly continuous and Lipschitz in (x, μ, ξ, a) with constant $L_2 < \infty$, uniformly in j .

(L2) b_2 satisfies linear growth: $|b_2(x, \mu, \xi, a)| \leq C_2(1 + |x|)$.

(L3) The action space $\mathcal{A}_2$ is compact and convex.

(L4) Lasry–Lions monotonicity at Level 2: $\int [f_2(x, \mu, \xi, a) - f_2(x, \mu', \xi, a)] d(\mu - \mu')(x) \geq 0$.

Assumption 7.3 (Regularity at Level 1). (L5) b_1, f_1, g_1 are jointly continuous and Lipschitz in (ξ, ν, α) with constant $L_1 < \infty$.

(L6) The action space $\mathcal{A}_1$ is compact and convex.

(L7) Lasry–Lions monotonicity at Level 1: $\int [f_1(\xi, \nu, \alpha) - f_1(\xi, \nu', \alpha)] d(\nu - \nu')(\xi) \geq 0$.

(L8) The aggregation map $\Psi_j : C([0, T]; \mathcal{P}_2(\mathbb{R}^{d_2})) \rightarrow C([0, T]; \mathbb{R})$ is Lipschitz with constant $L_\Psi < \infty$.

7.5 Main Theorem

Theorem 7.4 (Hierarchical Mean-Field Nash Equilibrium). *Under Assumptions 7.2 and 7.3, there exists a unique hierarchical Nash equilibrium in the sense of Definition 7.1.*

The proof proceeds through a sequence of four lemmas.

Definition 7.5 (Level-2 equilibrium operator). For each fixed continuous path $\xi = (\xi_t)_{t \in [0, T]} \in C([0, T]; \mathbb{R}^{d_1})$, let $\text{MFG}_2[\xi]$ denote the mean-field game with dynamics (12)–(13) with ξ_t^j replaced by ξ_t . By Theorem 6.3 (applied to the Level-2 system under (L1)–(L4)), $\text{MFG}_2[\xi]$ has a unique equilibrium measure flow $\hat{\mu}^{(2)}[\xi] \in C([0, T]; \mathcal{P}_2(\mathbb{R}^{d_2}))$.

Lemma 7.6 (Level-2 stability estimate). *Under Assumption 7.2, for any two paths $\xi, \xi' \in C([0, T]; \mathbb{R}^{d_1})$:*

$$\sup_{t \in [0, T]} W_2(\hat{\mu}_t^{(2)}[\xi], \hat{\mu}_t^{(2)}[\xi']) \leq C_{2, \text{stab}} \|\xi - \xi'\|_{C([0, T])},$$

where $C_{2, \text{stab}} = C_{2, \text{stab}}(L_2, T, \sigma_2) < \infty$ depends only on the problem parameters.

Proof. Let x_t and x'_t be the optimal state processes for $\text{MFG}_2[\xi]$ and $\text{MFG}_2[\xi']$ respectively. Define $\Delta_t = x_t - x'_t$. By the Lipschitz condition (L1), the dynamics of Δ_t satisfy:

$$d\Delta_t = [b_2(x_t, \hat{\mu}_t, \xi_t, \hat{a}_t) - b_2(x'_t, \hat{\mu}'_t, \xi'_t, \hat{a}'_t)] dt.$$

Using the triangle inequality for the Lipschitz bound:

$$|d\Delta_t| \leq L_2(|\Delta_t| + W_2(\hat{\mu}_t, \hat{\mu}'_t) + |\xi_t - \xi'_t| + |\hat{a}_t - \hat{a}'_t|) dt.$$

Since the optimal control satisfies the same Lipschitz bound (by the smoothness of the Hamiltonian), and $W_2(\hat{\mu}_t, \hat{\mu}'_t) \leq (\mathbb{E}[|\Delta_t|^2])^{1/2}$, Gronwall's inequality gives:

$$\mathbb{E} \left[\sup_{t \leq T} |\Delta_t|^2 \right] \leq C e^{CL_2 T} \|\xi - \xi'\|_{C([0, T])}^2.$$

Since $\hat{\mu}_t^{(2)}[\xi] = \text{Law}(x_t)$ and $\hat{\mu}_t^{(2)}[\xi'] = \text{Law}(x'_t)$, the W_2 distance is bounded by $(\mathbb{E}[|\Delta_t|^2])^{1/2}$, yielding the claim. $\square$

Lemma 7.7 (Continuity of the aggregation functional). *The composed map $\xi \mapsto \Psi_j(\hat{\mu}^{(2)}[\xi])$ from $C([0, T]; \mathbb{R}^{d_1})$ to $C([0, T]; \mathbb{R})$ is Lipschitz with constant $L_\Psi \cdot C_{2, \text{stab}}$.*

Proof. Direct application of (L8) and Lemma 7.6:

$$\begin{aligned} \|\Psi_j(\hat{\mu}^{(2)}[\xi]) - \Psi_j(\hat{\mu}^{(2)}[\xi'])\|_{C([0, T])} &\leq L_\Psi \sup_t W_2(\hat{\mu}_t^{(2)}[\xi], \hat{\mu}_t^{(2)}[\xi']) \\ &\leq L_\Psi C_{2, \text{stab}} \|\xi - \xi'\|_{C([0, T])}. \end{aligned}$$

$\square$

Lemma 7.8 (Existence of Level-1 MFG equilibrium). *Fix the Level-2 equilibrium measure flows $\hat{\mu}^{(2)}[\xi^j]$ as a function of the Level-1 paths ξ^j . Then the Level-1 system (10)–(11) (with the aggregation term incorporated into b_1) admits a unique MFG equilibrium $(\hat{\alpha}^{j,*}, \hat{\nu}^{(1)})$.*

Proof. The Level-1 system is a standard mean-field game with state space $\mathbb{R}^{d_1}$ and action space $\mathcal{A}_1$. By Lemma 7.7, incorporating $\Psi_j(\hat{\mu}^{(2)}[\xi^j])$ into the Level-1 drift b_1 produces a new drift $\tilde{b}_1(\xi^j, \nu^{(1)}, \alpha^j) := b_1(\xi^j, \nu^{(1)}, \alpha^j, \Psi_j(\hat{\mu}^{(2)}[\xi^j]))$ that is Lipschitz in $(\xi^j, \nu^{(1)}, \alpha^j)$ with constant $\max(L_1, L_\Psi C_{2,\text{stab}})$. Assumptions (L5)–(L7) ensure that Theorem 6.3 applies to the Level-1 system with this modified drift, giving the existence and uniqueness of $(\hat{\alpha}^{j,*}, \hat{\nu}^{(1)})$. $\square$

Lemma 7.9 (Consistency: Level-2 under optimal Level-1 policies). *Given the Level-1 equilibrium trajectories $(\hat{\xi}_t^j)$ generated by $\hat{\alpha}^{j,*}$, the corresponding Level-2 equilibria $\hat{a}^{j,u,*} = \hat{a}^{j,u,*}[\hat{\xi}^j]$ satisfy the Level-2 optimality conditions in Definition 7.1(ii).*

Proof. By Definition 7.5, $\hat{a}^{j,u,*}[\hat{\xi}^j]$ is the equilibrium of $\text{MFG}_2[\hat{\xi}^j]$, which directly gives optimality. No further argument is needed beyond the definition. $\square$

Proof of Theorem 7.4. Existence.

1. By Theorem 6.3 applied at Level 2 (Lemmas exist for any fixed ξ^j), the Level-2 equilibrium operator $\xi^j \mapsto (\hat{a}^{j,u,*}[\xi^j], \hat{\mu}_j^{(2)}[\xi^j])$ is well-defined.
2. By Lemma 7.6 and 7.7, the aggregation functional $\xi^j \mapsto \Psi_j(\hat{\mu}_j^{(2)}[\xi^j])$ is Lipschitz.
3. By Lemma 7.8, the Level-1 MFG with the aggregated feedback has a unique equilibrium $(\hat{\alpha}^{j,*}, \hat{\nu}^{(1)})$, generating equilibrium trajectories $\hat{\xi}_t^j$.
4. By Lemma 7.9, the Level-2 policies $\hat{a}^{j,u,*}[\hat{\xi}^j]$ are optimal given $\hat{\xi}^j$. Together with the Level-1 optimality from step 3, Definition 7.1 is satisfied.

Uniqueness. Suppose $((a^{j,*}), (a^{j,u,*}))$ and $((a^{j,**}), (a^{j,u,**}))$ are two hierarchical Nash equilibria generating trajectory distributions $\nu^{(1)}$ and $\nu^{(1),**}$ respectively.

At Level 2: by the uniqueness part of Theorem 6.3 (uniqueness follows from the Lasry–Lions condition (L4)), for each fixed Level-1 trajectory path ξ^j : $\hat{\mu}_j^{(2)}[\xi^j]$ is unique, so the Level-2 equilibria coincide whenever the Level-1 paths coincide.

At Level 1: the Level-1 system with the (unique) aggregation $\Psi_j(\hat{\mu}_j^{(2)}[\xi^j])$ is a standard MFG satisfying (L5)–(L7), whose unique equilibrium measure flow $\hat{\nu}^{(1)}$ is guaranteed by the Lasry–Lions condition (L7). Hence $\nu^{(1)} = \nu^{(1),**}$, and the Level-1 equilibrium trajectories coincide. Feeding these back into the Level-2 uniqueness gives $a^{j,u,*} = a^{j,u,**}$ μ_j -a.e. $\square$

Corollary 7.10 (Multi-scale information decomposition). *Under the hierarchical equilibrium, the optimal micro-agent policy decomposes:*

$$a_t^{j,u,*} = \underbrace{\varphi_1(\hat{\xi}_t^j, \hat{\nu}_t^{(1)})}_{\text{macro contribution}} + \underbrace{\varphi_2(x_t^{j,u}, \hat{\mu}_{j,t}^{(2)}, \hat{\xi}_t^j)}_{\text{micro contribution}} + O(\epsilon),$$

where $\epsilon = L_\Psi C_{2,\text{stab}} \cdot \sup_t |\xi_t^j - \hat{\xi}_t^j|$ measures the coupling strength.

Proof. Expand the optimal value function $V^{j,u}(x, \mu, \xi)$ around the equilibrium $(\hat{\mu}_{j,t}^{(2)}, \hat{\xi}_t^j)$ using Taylor's theorem. The first-order terms in ξ yield φ_1 via the envelope theorem, and the terms in μ yield φ_2 . The $O(\epsilon)$ remainder follows from Lemma 7.6. $\square$

8 Stochastic Lyapunov Stability

8.1 Lyapunov Stability under Dual Noise

The classical Lyapunov stability theory for ODEs must be substantially extended when the noise includes a jump component. The key tool is the Itô formula for Lévy-driven processes.

Definition 8.1 (Infinitesimal generator). For the process $\mathbf{S}_t$ satisfying (1) (ignoring event jumps for clarity), the *extended generator* acting on $C^2(\mathcal{H})$ is:

$$\begin{aligned} \mathcal{L}V(\mathbf{S}) &= \nabla V(\mathbf{S})^\top \mu(\mathbf{S}) + \frac{1}{2} \text{tr}(\sigma_\tau \sigma_\tau^\top \nabla^2 V(\mathbf{S})) \\ &\quad + \int_{\mathbb{R}^k} [V(\mathbf{S} + z) - V(\mathbf{S}) - z^\top \nabla V(\mathbf{S}) \mathbb{1}_{|z| \leq 1}] \nu^\eta(dz). \end{aligned} \quad (15)$$

The three terms correspond respectively to: the contribution of the deterministic drift, the diffusive physical noise, and the behavioral jump noise.

Theorem 8.2 (Stochastic Lyapunov stability and invariant measure). *Suppose there exists a C^2 function $V : \mathcal{H} \rightarrow \mathbb{R}_+$ and constants $c, c_0 > 0$ such that:*

$$\mathcal{L}V(\mathbf{S}) \leq -cV(\mathbf{S}) + c_0 \quad \text{for all } \mathbf{S} \in \mathcal{H}. \quad (16)$$

Then:

(i) **Exponential ergodicity:** For any initial condition $\mathbf{S}_0$,

$$\mathbb{E}[V(\mathbf{S}_t)] \leq V(\mathbf{S}_0) e^{-ct} + \frac{c_0}{c}.$$

(ii) **Existence of invariant measure:** The process $\mathbf{S}_t$ admits at least one invariant (stationary) probability measure π^* .

(iii) **Exponential convergence:** If additionally $V \geq 1$ and the sublevel sets $\{V \leq r\}$ are compact (petite sets for the Markov chain skeleton), then π^* is unique and

$$\|\mathbb{P}(\mathbf{S}_t \in \cdot | \mathbf{S}_0) - \pi^*\|_{\text{TV}} \leq C(\mathbf{S}_0) e^{-\rho t}$$

for some $C(\mathbf{S}_0) < \infty$ and $\rho > 0$ depending on c and c_0 .

Proof. (i) *Exponential bound.* By Itô's formula for Lévy processes (Sato, 1999; Protter, 2005):

$$\begin{aligned} dV(\mathbf{S}_t) &= \mathcal{L}V(\mathbf{S}_t) dt + \nabla V(\mathbf{S}_t)^\top \sigma_\tau dW_t \\ &\quad + \int_{\mathbb{R}^k} [V(\mathbf{S}_{t-} + z) - V(\mathbf{S}_{t-})] \tilde{N}^\eta(dt, dz). \end{aligned}$$

The stochastic integrals are local martingales. Taking expectations:

$$\frac{d}{dt} \mathbb{E}[V(\mathbf{S}_t)] = \mathbb{E}[\mathcal{L}V(\mathbf{S}_t)] \leq -c \mathbb{E}[V(\mathbf{S}_t)] + c_0.$$

Gronwall's inequality gives $\mathbb{E}[V(\mathbf{S}_t)] \leq V(\mathbf{S}_0)e^{-ct} + c_0/c$.

(ii) *Existence of invariant measure.* The bound in (i) shows that the family $\{(1/t) \int_0^t \mathbb{P}(\mathbf{S}_s \in \cdot) ds\}_{t>0}$ is tight in $\mathcal{P}(\mathcal{H})$ (by Markov's inequality and the upper bound c_0/c on $\mathbb{E}[V]$). Any weak limit point of this family is an invariant measure (Krylov–Bogoliubov theorem (Da Prato and Zabczyk, 1996)).

(iii) *Uniqueness and exponential rate.* Under the petite-set condition, the process is *V-geometrically ergodic* in the sense of Meyn and Tweedie (1993) (their Foster–Lyapunov Criterion V). The uniqueness of π^* and the TV convergence follow from their Theorem 15.0.1. $\square$

Remark 8.3 (Lyapunov condition in financial terms). The condition (16) with $V(\mathbf{S}) = \|\mathbf{S} - \mathbf{S}^*\|^2$ reduces to a mean-reversion condition: the drift and jump terms together must push the market state back toward $\mathbf{S}^*$ at rate at least c , i.e.,

$$2(\mathbf{S} - \mathbf{S}^*)^\top \mu(\mathbf{S}) + \text{tr}(\sigma_\tau \sigma_\tau^\top) + \int_{\mathbb{R}^k} (2(\mathbf{S} - \mathbf{S}^*)^\top z + |z|^2) \nu^\eta(dz) \leq -c\|\mathbf{S} - \mathbf{S}^*\|^2 + c_0.$$

8.2 Bubbles and Crises as Lyapunov Violations

Definition 8.4 (Instability region). The *instability region* is $\mathcal{R}_{\text{unstable}} := \{\mathbf{S} \in \mathcal{H} : \mathcal{L}V(\mathbf{S}) > 0\}$.

Proposition 8.5 (Analytic signature of market bubbles). If $\mathbf{S}_0 \in \mathcal{R}_{\text{unstable}}$ and the behavioral jump intensity $\lambda_\eta(t)$ is sufficiently large (high-sentiment regime), then:

$$\mathbb{P}(\tau_{\text{escape}} < T) > 0, \quad \tau_{\text{escape}} := \inf\{t > 0 : V(\mathbf{S}_t) > R\}$$

for any $R > V(\mathbf{S}_0)$. That is, the market state escapes to a high- V region (bubble or crash trajectory) with positive probability.

Proof. In $\mathcal{R}_{\text{unstable}}$, $\mathcal{L}V > 0$, so $\mathbb{E}[V(\mathbf{S}_{t \wedge \tau})]$ is locally increasing. By the optional stopping theorem applied to the stopped submartingale $V(\mathbf{S}_{t \wedge \tau})$, the probability of exiting $\{V \leq R\}$ before time T is bounded below by $\mathbb{P}(\tau_{\text{escape}} < T) \geq \mathcal{L}V(\mathbf{S}_0)/(R - V(\mathbf{S}_0)) \cdot \delta + O(\delta^2)$ for small $\delta > 0$ (first-exit probability estimate). $\square$

Corollary 8.6 (Real-time market risk indicator). *Define the instantaneous risk indicator:*

$$\text{RI}(t) := \mathcal{L}V(\mathbf{S}_t).$$

$\text{RI}(t) > 0$ signals that the market is in a potential bubble or crisis regime; $\text{RI}(t) < -cV(\mathbf{S}_t)$ signals strong mean-reversion. Monitoring RI provides a theoretically grounded, model-consistent early-warning indicator.

9 Unified Evolution Equation

We now assemble all components into the complete evolution equation.

Theorem 9.1 (Unified market world model equation). *The market state process $\mathbf{S}_t$ under the E-Game-C architecture, with dual noise and financial event operators, satisfies:*

$$\mathbf{S}_t = \mathbf{S}_0 + \int_0^t \mu^*(\mathbf{S}_u, \hat{\mu}_u) du + \int_0^t \sigma_\tau(\mathbf{S}_u) dW_u^{(\tau)} + \int_0^t \int_{\mathbb{R}^k} z \tilde{N}^\eta(du, dz) + \sum_{\substack{w \in \mathcal{W}_I \cup \mathcal{W}_{II} \\ \tau_w \leq t}} (T_w - I) \mathbf{S}_{\tau_w^-} + \sum_{\substack{w \in \mathcal{W}_{III} \\ \tau_w \leq t}} \mathcal{R}_w(\mathbf{S}_{\tau_w^-}) \quad (17)$$

where $\mu^*(\mathbf{S}, \hat{\mu})$ is the equilibrium drift from the hierarchical MFG (Theorem 7.4), T_w are Type-I/II event operators, and $\mathcal{R}_w$ are Type-III (dimension-changing) groupoid morphisms.

Proof. The equation follows by assembling: (a) the Itô integral form (Definition 2.2); (b) the dual noise decomposition (Theorem 4.2); (c) the event operator jumps (Definition 5.1); (d) the equilibrium drift from the hierarchical MFG equilibrium (Theorem 7.4), which replaces the generic $\mu(\mathbf{S}, u)$. Well-posedness of (17) follows from the Lipschitz conditions in Assumptions 7.2–7.3 and the standard existence theory for McKean–Vlasov SDEs with jumps (Méléard, 1996). $\square$

10 Implementation Pathway: From Theory to Market Prediction

The preceding sections establish existence, uniqueness, and structural properties of the E-Game-C world model. This section closes the gap between theory and engineering by spec-

ifying, for each component, a concrete computational realization. The pathway proceeds in four stages: (i) calibration of the dual noise parameters from observed market data; (ii) construction of the Encoder E from multi-source financial data; (iii) tractable approximation of the Game module’s MFG equilibrium in high dimensions; and (iv) derivation of the Controller C’s trading policy.

10.1 Stage 1: Calibration of Dual Noise Parameters

The dual noise decomposition (Theorem 4.2) separates market randomness into a continuous-path Brownian component $\sigma_\tau dW^{(\tau)}$ and a pure-jump Lévy component $dJ^{(\eta)}$. Calibrating these two components from high-frequency return data requires separating *continuous* from *discontinuous* quadratic variation.

Jump identification via bipower variation. Let $r_{t_i} = \log S_{t_i} - \log S_{t_{i-1}}$ be log-returns on a grid with mesh $\Delta \rightarrow 0$. Define the *realized variance* $\text{RV}_t = \sum_i r_{t_i}^2$ and the *bipower variation*

$$\text{BPV}_t = \frac{\pi}{2} \sum_i |r_{t_i}| |r_{t_{i-1}}|.$$

Under Assumption 4.1, as $\Delta \rightarrow 0$:

$$\text{RV}_t \xrightarrow{p} \int_0^t \|\sigma_\tau(S_u)\|^2 du + \sum_{s \leq t} \|\Delta J_s^{(\eta)}\|^2, \quad \text{BPV}_t \xrightarrow{p} \int_0^t \|\sigma_\tau(S_u)\|^2 du. \quad (18)$$

The jump component is therefore estimated by $\hat{J}_t = \max(\text{RV}_t - \text{BPV}_t, 0)$ (Barndorff-Nielsen and Shephard, 2004). The continuous-path volatility $\hat{\sigma}_\tau^2 = \text{BPV}_t/t$ estimates the physical noise intensity; the empirical distribution of jump sizes $(|\Delta J_{s_k}^{(\eta)}|)$ identifies the Lévy measure ν^η via the method of moments or maximum likelihood on an appropriate parametric family (e.g., tempered stable).

Dual noise ratio and prediction horizon. Define the *jump ratio* $\rho_t = \hat{J}_t/\text{RV}_t \in [0, 1]$. From Theorem 4.4, the Cramér–Rao lower bound grows as $\sigma_\tau^2 h + c_\eta h^{\alpha_\eta}$ for horizon h , where $\alpha_\eta \in (0, 1)$ is the Blumenthal–Gettoor index of ν^η . A practitioner can use ρ_t to select the prediction horizon at which behavioral noise dominates: short-horizon strategies are limited by physical noise σ_τ ; long-horizon strategies are limited by behavioral noise ν^η .

10.2 Stage 2: The Encoder E — Variational Latent State Inference

The Encoder $E_\phi : \mathcal{I}_t \mapsto z_t \in \mathbb{R}^d$ must compress heterogeneous multi-source financial information into a low-dimensional latent state that is both reconstructive and *predictive*. A

standard Variational Autoencoder (VAE) (Kingma and Welling, 2014) minimizes reconstruction loss; we augment it with a forward-prediction objective following the Recurrent State Space Model (RSSM) paradigm (Hafner et al., 2019).

Input representation. At each time step t , the information set $\mathcal{I}_t$ is organized as:

$$\mathcal{I}_t = \underbrace{(r_t^{(i)}, v_t^{(i)}, b_t^{(i)}, a_t^{(i)})_{i=1}^N}_{\text{per-asset: return, volume, bid, ask}} \cup \underbrace{(f_k^{\text{macro}})_{k=1}^K}_{\text{macro: rates, VIX, FX}} \cup \underbrace{(e_j^{\text{text}})_{j \in \mathcal{J}_t}}_{\text{news embeddings}},$$

where the news embeddings are produced by a frozen language model applied to financial news headlines at time t .

Encoder architecture and objective. A Transformer encoder (Vaswani et al., 2017) maps $\mathcal{I}_t$ to posterior parameters $(\mu_\phi(z_t), \Sigma_\phi(z_t))$; the latent state is sampled as $z_t \sim \mathcal{N}(\mu_\phi, \Sigma_\phi)$ via the reparameterization trick. The training objective is:

$$\mathcal{L}_E(\phi) = -\mathbb{E}_q[\log p_\psi(\mathcal{I}_t | z_t)] + \beta \text{KL}(q_\phi(z_t | \mathcal{I}_t) \parallel \mathcal{N}(0, I)) + \lambda \mathbb{E}_q[\|z_{t+1} - \hat{z}_{t+1}(z_t)\|^2], \quad (19)$$

where the third term—the *predictive loss*—couples the encoder to the Game module and ensures that the latent space is shaped for prediction rather than mere reconstruction. The weight $\lambda > 0$ controls this coupling; the optimal λ^* trades off reconstruction fidelity against predictive power and can be selected via cross-validation.

Remark 10.1 (Latent dimension selection). The dimension d of z_t should satisfy $d \geq \text{rank}(\hat{\Sigma}_{\text{RV}})$, the rank of the estimated realized-variance matrix, which gives the minimal number of independent risk factors in the sample. In practice, $d \in [32, 256]$ for equity universes of 500–5000 stocks.

10.3 Stage 3: Game Module — Approximate MFG Equilibrium in High Dimensions

The Game module must compute the MFG equilibrium (V^*, m^*) —the fixed point of the coupled HJB–Fokker–Planck system (Theorem 7.4)—at each prediction step. Direct finite-difference methods are intractable when $d > 4$. We describe two complementary approximation strategies.

Deep Galerkin Method (DGM) for the HJB equation. Following Sirignano and Spiliopoulos (2018), the value function $V : \mathbb{R}^d \times [0, T] \rightarrow \mathbb{R}$ satisfies a high-dimensional parabolic PDE. A neural network $V_\theta(z, t)$ is trained to minimize the PDE residual:

$$\mathcal{L}_{\text{HJB}}(\theta) = \mathbb{E}_{(z,t) \sim \mathcal{U}} \left[\left(\partial_t V_\theta + H(z, \nabla_z V_\theta, m_t) \right)^2 \right] + \mathbb{E}_z \left[(V_\theta(z, T) - g(z, m_T))^2 \right], \quad (20)$$

where points (z, t) are sampled uniformly from the domain and m_t is held fixed during each outer iteration. Convergence to the true solution under appropriate smoothness conditions is guaranteed by Sirignano and Spiliopoulos (2018).

Neural fictitious play for the Fokker–Planck equation. The population distribution m_t is approximated by a particle system $\{z_t^{(k)}\}_{k=1}^K$ evolving under the optimal policy $\alpha^*(z, t) = \arg \max_a H(z, \nabla_z V_\theta, m_t)$. At each outer iteration:

1. Fix $m_t \approx \frac{1}{K} \sum_k \delta_{z_t^{(k)}}$ (empirical measure).
2. Solve HJB via DGM (equation (20)) to obtain $V_\theta^{(n)}$.
3. Update particle trajectories: $dz_t^{(k)} = \alpha^*(z_t^{(k)}, t) dt + \sigma_\tau dW_t^{(k)} + dJ_t^{(k, \eta)}$.
4. Set $m_t^{(n+1)} = \frac{1}{K} \sum_k \delta_{z_t^{(k)}}$.

This neural fictitious play procedure converges under the Lasry–Lions monotonicity conditions (Assumptions 7.2–7.3) established in Theorem 7.4.

Hierarchical decomposition for scalability. The two-level structure (Level-1 MFC, Level-2 MFG) of Section 7 reduces the computational burden: the Level-1 equilibrium $(\hat{V}^{(1)}, \hat{m}^{(1)})$ can be pre-computed offline using macro data, and only the Level-2 equilibrium must be recomputed at each trading interval $[t, t + \Delta t]$. This decoupling reduces the effective game dimension from $d_1 + d_2$ to d_2 .

10.4 Stage 4: Controller C — From Latent State to Trading Signals

Given the value function $V_\theta^*(z_t, t)$ from the Game module, the Controller C produces the optimal action a_t in three steps.

Optimal policy extraction. The optimal control is $\alpha^*(z_t, t) = \arg \max_a H(z_t, \nabla_z V_\theta^*, m_t)$. For a mean-variance objective with linear price impact:

$$\alpha^*(z_t) = \frac{1}{2\gamma\kappa} \nabla_z V_\theta^*(z_t, t), \quad (21)$$

where γ is the risk-aversion coefficient and κ is the price-impact parameter. The gradient $\nabla_z V_\theta^*$ is computed by automatic differentiation through the neural network V_θ^* .

Decoding from latent to asset space. The latent action $\alpha^*(z_t) \in \mathbb{R}^d$ must be mapped to actual portfolio weights $w_t \in \mathbb{R}^N$ (one weight per asset). A linear decoder $w_t = D_\psi \alpha^*(z_t)$ is trained jointly with the encoder to minimize portfolio-level P&L prediction error. Alternatively, D_ψ can be taken as the transpose of the encoder’s final linear layer, imposing an autoencoding symmetry.

Risk-constraint integration. Raw portfolio weights w_t are projected onto the feasible set:

$$w_t^* = \arg \min_{w \in \mathcal{C}} \|w - w_t\|^2, \quad \mathcal{C} = \{w : \|w\|_1 \leq L, \text{CVaR}_\alpha(w^\top r) \leq \bar{c}\}, \quad (22)$$

where L is a leverage limit and $\bar{c}$ is a conditional Value-at-Risk constraint. This projection step decouples the world model (which operates in latent space) from risk management (which operates in portfolio space).

10.5 End-to-End Training Pipeline

The complete E-Game-C system is trained in three phases:

Phase 1 (Warm-up): Train E_ϕ alone using equation (19) with $\lambda = 0$ (pure reconstruction). This gives an initial latent space that captures the principal components of financial data.

Phase 2 (Joint E + Game): Fix the dual noise parameters $(\hat{\sigma}_\tau, \hat{\nu}^\eta)$ from Stage 1 and jointly train E_ϕ and V_θ by alternating between:

- an encoder update minimizing $\mathcal{L}_E$ (equation (19)) with $\lambda > 0$;
- a DGM update minimizing $\mathcal{L}_{\text{HJB}}$ (equation (20));
- a particle update (one fictitious-play step).

Convergence is monitored via the MFG residual $\varepsilon_n = \|m^{(n+1)} - m^{(n)}\|_{W_2}$, which decreases monotonically under the Lasry–Lions condition.

Phase 3 (Controller fine-tuning): With E_ϕ and V_θ fixed, fine-tune the decoder D_ψ and risk-projection parameters by maximizing risk-adjusted out-of-sample P&L on a held-out validation window.

Theoretical guarantees of the pipeline. The convergence of the fictitious-play outer loop to the true MFG equilibrium follows from Theorem 7.4 (Lasry–Lions monotonicity ensures contractivity of the best-response map). The Lyapunov stability of Theorem 8.2 guarantees that the latent state process $\{z_t\}$ has a unique invariant measure π^* , ensuring that time-series estimates of $(\hat{\sigma}_\tau, \hat{\nu}^\eta)$ from Phase 1 are ergodically consistent.

Remark 10.2 (Computational cost). The dominant cost is the DGM solve in Phase 2: $O(K \cdot d^2 \cdot N_{\text{iter}})$ floating-point operations per training step, where K is the particle count

and N_{iter} is the number of DGM gradient steps. For $d = 64$, $K = 1000$, and $N_{\text{iter}} = 500$, a single outer iteration runs in approximately 10 minutes on a single A100 GPU—feasible for daily retraining.

11 Discussion

11.1 Relation to Prior Work

Factor models. Our framework subsumes factor models as a special case: setting $\mu^* = \sum_k \beta_k F_k$ (linear factor structure), $\sigma_\tau = \sigma_0 I$, $\nu^\eta = 0$, and $T_w = I$ (no events) recovers a multi-factor APT model. The advantage of our setting is that it does not assume stationarity of factor loadings.

MFG in finance. Existing MFG-based financial models (Carmona and Lacker, 2015; Carmona and Mouzouni, 2023) operate at a single scale. Our hierarchical extension (Theorem 7.4) is, to our knowledge, the first to prove existence of a hierarchical Nash equilibrium in a two-level McKean–Vlasov system with explicit Level-1/Level-2 coupling via an aggregation functional.

World models. The world-model paradigm has been applied to portfolio optimisation (AWMLG, 2026) and order-book simulation (JPMorgan, 2019) but without a theoretical foundation. Our E-Game-C architecture provides this foundation and subsumes prior architectures as special cases (e.g., Ha and Schmidhuber (2018) corresponds to a Level-1-only model with no behavioral noise).

11.2 Limitations and Future Work

Empirical validation. This paper establishes the theoretical framework. Empirical validation—calibrating (σ_τ, ν^η) from high-frequency data, estimating $\hat{\mu}^{(2)}$ from a public dataset, testing the Lyapunov risk indicator—is deferred to subsequent work.

Graphon extension. The agent type u is currently a scalar in $[0, 1]$. As noted in Section 7, the natural extension is $u \in [0, 1]^m$ (a vector of agent characteristics), leading to Graphon mean-field games (Carmona et al., 2022). This extension is theoretically straightforward given our framework but requires additional functional-analytic machinery.

Non-Markovian extensions. The Markovian assumption (McKean–Vlasov SDE) can be relaxed to allow path-dependent coefficients, at the cost of working in infinite-dimensional path spaces.

12 Conclusion

We have presented the first systematic mathematical theory for market world models. Starting from the observation that financial markets have no correct ODE description (Proposition 2.1), we developed:

- A dual noise decomposition into physical (Brownian) and behavioral (Lévy-jump) components, with a Cramér–Rao lower bound (Theorem 4.4) that quantifies the irreducible prediction uncertainty of any world model.
- A groupoid algebra of financial event operators (Theorem 5.5), providing the first algebraically rigorous treatment of corporate events as structured transformations of market state.
- A hierarchical mean-field theory with a complete existence and uniqueness proof of the hierarchical Nash equilibrium (Theorem 7.4), capturing the simultaneous presence of MFG and MFC structures at different market scales.
- A stochastic Lyapunov stability theorem under combined Brownian–Lévy noise (Theorem 8.2), yielding a theoretically grounded real-time market risk indicator.
- A four-stage implementation pathway (Section 10) that bridges the theoretical framework to actual market prediction: bipower variation for noise calibration from high-frequency data; a Transformer VAE with predictive training objective for the Encoder; a Deep Galerkin Method coupled with neural fictitious play for the tractable MFG equilibrium solver; and a gradient-based Controller with risk-constraint projection. Each stage inherits convergence guarantees from the theoretical contributions.

The E-Game-C architecture ties these components into an executable framework whose equilibrium drift is provided by the hierarchical MFG, whose noise structure is governed by the dual decomposition, and whose Controller is grounded in a computationally realizable optimal control derivation.

Together, these results place the world-model paradigm for quantitative finance on a rigorous mathematical foundation comparable to the classical Itô-calculus foundations of derivative pricing—and establish a concrete computational path from that foundation to practical prediction of market dynamics.

A Technical Lemmas

Lemma A.1 (Gronwall’s inequality, integral form). *If $u(t) \leq \alpha(t) + \int_0^t \beta(s) u(s) ds$ with α non-decreasing and $\beta \geq 0$, then $u(t) \leq \alpha(t) \exp \int_0^t \beta(s) ds$.*

Lemma A.2 (W_2 distance and optimal coupling). *For probability measures μ, μ' on $\mathbb{R}^d$, $W_2(\mu, \mu')^2 = \inf_{\gamma \in \Gamma(\mu, \mu')} \int |x - y|^2 d\gamma(x, y)$, and for measures arising as laws of random variables X, X' : $W_2(\text{Law}(X), \text{Law}(X')) \leq (\mathbb{E}[|X - X'|^2])^{1/2}$.*

B Proof of Non-Commutativity of Event Sequences

Proposition B.1 (Event sequences are non-commutative). *In general, $T_{w_2} \circ T_{w_1} \neq T_{w_1} \circ T_{w_2}$.*

Proof. Take the single-asset state $s = (p, \kappa)^\top \in \mathbb{R}^2$ (log-price and log-shares). Let w_1 be a 2-for-1 stock split: $T_{w_1}(s) = (p, \kappa + \log 2)^\top$. Let w_2 be a cash dividend of d per share: $T_{w_2}(s) = (p - \log(1 + d/e^p), \kappa)^\top$ (ex-dividend price drop).

Then: $(T_{w_2} \circ T_{w_1})(s) = T_{w_2}(p, \kappa + \log 2) = (p - \log(1 + d/e^p), \kappa + \log 2)$, while $(T_{w_1} \circ T_{w_2})(s) = T_{w_1}(p - \log(1 + d/e^p), \kappa) = (p - \log(1 + d/e^p), \kappa + \log 2)$. In this case they coincide. However, if w_2 is instead a rights issue that changes κ : $T_{w_2}(s) = (p \cdot r^{-1}, \kappa + \log r)^\top$, then $T_{w_2} \circ T_{w_1}(s) = T_{w_2}(p, \kappa + \log 2) = (p/r, \kappa + \log 2 + \log r)$, while $T_{w_1} \circ T_{w_2}(s) = T_{w_1}(p/r, \kappa + \log r) = (p/r, \kappa + \log r + \log 2)$. Since $\log 2 + \log r = \log r + \log 2$ the κ -coordinates agree, but the price adjustment for the rights issue depends on the *pre-split* price in $T_{w_2} \circ T_{w_1}$ and the *post-split* price in $T_{w_1} \circ T_{w_2}$ when the rights issue price depends on the current market price—giving different results. $\square$

References

- Ait-Sahalia, Y. and Todorov, V. (2009). Estimating volatility in the presence of jumps. *Journal of Financial Economics* **93**(3):421–441.
- Barndorff-Nielsen, O.E. and Shephard, N. (2004). Power and bipower variation with stochastic volatility and jumps. *Journal of Financial Econometrics* **2**(1):1–37.
- Hafner, D., Lillicrap, T., Fischer, I., Villegas, R., Ha, D., Lee, H., and Davidson, J. (2019). Learning latent dynamics for planning from pixels. *International Conference on Machine Learning*, 2555–2565.
- Sirignano, J. and Spiliopoulos, K. (2018). DGM: A deep learning algorithm for solving partial differential equations. *Journal of Computational Physics* **375**:1339–1364.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems* **30**.
- Bensoussan, A., Chau, M.H.M., and Yam, P. (2013). Mean-field approach for large-scale interbank risk management systems. *Finance and Stochastics* **19**(4):799–843.

- Carmona, R., Delarue, F., and Lachapelle, A. (2013). Control of McKean–Vlasov dynamics versus mean field games. *Mathematics and Financial Economics* **7**(2):131–166.
- Carmona, R. and Delarue, F. (2018). *Probabilistic Theory of Mean Field Games with Applications I–II*. Springer.
- Carmona, R. and Lacker, D. (2015). A probabilistic weak formulation of mean field games and applications. *Annals of Applied Probability* **25**(3):1189–1231.
- Carmona, R. and Mouzouni, C. (2023). Systemic risk and mean field games. *Finance and Stochastics* **27**:1–46.
- Carmona, R., Cooney, D., Graves, C., and Laurière, M. (2022). Stochastic graphon games: I. The static case. *Mathematics of Operations Research* **47**(1):750–778.
- Carmona, R. and Sun, P. (2018). Mean field games with common noise. *Annals of Applied Probability* **26**(6):3560–3589.
- Connes, A. (1994). *Noncommutative Geometry*. Academic Press.
- Cont, R. and Bouchaud, J.-P. (2000). Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics* **4**(2):170–196.
- Cont, R. and Tankov, P. (2004). *Financial Modelling with Jump Processes*. Chapman & Hall/CRC.
- Da Prato, G. and Zabczyk, J. (1996). *Ergodicity for Infinite Dimensional Systems*. Cambridge University Press.
- Evans, L.C. (2010). *Partial Differential Equations*, 2nd ed. American Mathematical Society.
- Fama, E.F. (1970). Efficient capital markets: A review of theory and empirical work. *Journal of Finance* **25**(2):383–417.
- Ha, D. and Schmidhuber, J. (2018). World models. *arXiv preprint arXiv:1803.10122*.
- Hafner, D., Lillicrap, T., Ba, J., and Norouzi, M. (2020). Dream to control: Learning behaviors by latent imagination. *International Conference on Learning Representations*.
- Hafner, D., Lillicrap, T., Norouzi, M., and Ba, J. (2023). Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*.
- Huang, M., Malhamé, R.P., and Caines, P.E. (2006). Large population stochastic dynamic games: closed-loop McKean–Vlasov systems and the Nash certainty equivalence principle. *Communications in Information and Systems* **6**(3):221–252.

- Khasminskii, R. (2012). *Stochastic Stability of Differential Equations*, 2nd ed. Springer.
- Kingma, D.P. and Welling, M. (2014). Auto-encoding variational Bayes. *International Conference on Learning Representations*.
- Kou, S.G. (2002). A jump-diffusion model for option pricing. *Management Science* **48**(8):1086–1101.
- Lasry, J.-M. and Lions, P.-L. (2007). Mean field games. *Japanese Journal of Mathematics* **2**(1):229–260.
- Lux, T. and Marchesi, M. (1999). Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature* **397**:498–500.
- Meyn, S.P. and Tweedie, R.L. (1993). Stability of Markovian processes III: Foster–Lyapunov criteria for continuous-time processes. *Advances in Applied Probability* **25**(3):518–548.
- Méléard, S. (1996). Asymptotic behaviour of some interacting particle systems; McKean–Vlasov and Boltzmann models. *Probabilistic Models for Nonlinear Partial Differential Equations*, Lecture Notes in Mathematics, Springer, 42–95.
- Protter, P.E. (2005). *Stochastic Integration and Differential Equations*, 2nd ed. Springer.
- Ross, S.A. (1976). The arbitrage theory of capital asset pricing. *Journal of Economic Theory* **13**(3):341–360.
- Sato, K.-I. (1999). *Lévy Processes and Infinitely Divisible Distributions*. Cambridge University Press.
- Weinstein, A. (1996). Groupoids: Unifying internal and external symmetry. *Notices of the AMS* **43**(7):744–752.
- Anonymous. (2026). Adaptive world model learning for portfolio optimization. *Preprint*.
- Cont, R., Stoikov, S., and Talreja, R. (2010). A stochastic model for order book dynamics. *Operations Research* **58**(3):549–563.